

## #Stock Trading Using Deep Q-Learning

### Problem Statement

Prepare an agent by implementing Deep Q-Learning that can perform unsupervised trading in stock trade. The aim of this project is to train an agent that uses Q-learning and neural networks to predict the profit or loss by building a model and implementing it on a dataset that is available for evaluation.

The stock trading index environment provides the agent with a set of actions:

- Buy
- Sell
- Sit

This project has following sections:

- Import libraries
- Create a DQN agent
- Preprocess the data
- Train and build the model
- Evaluate the model and agent

### Steps to perform

In the section **create a DQN agent**, create a class called agent where:

- Action size is defined as 3
- Experience replay memory to deque is 1000
- Empty list for stocks that has already been bought
- The agent must possess the following hyperparameters:
  - $\gamma = 0.95$
  - $\epsilon = 1.0$
  - $\epsilon_{\text{final}} = 0.01$
  - $\epsilon_{\text{decay}} = 0.995$

Note: It is advised to compare the results using different values in hyperparameters.

- Neural network has 3 hidden layers
- Action and experience replay are defined

## Solution

### Dataset

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

```
!unzip -qq
/content/drive/MyDrive/datasets/simplilearn_RL_stock_trading/dataset.zip
```

```
!scp /content/dataset/GSPC_Training_Dataset.csv
/content/GSPC_Training_Dataset.csv
!scp /content/dataset/GSPC_Evaluation_Dataset.csv
/content/GSPC_Evaluation_Dataset.csv
```

### REINFORCEMENT LIBRARIES

```
!pip install python-opengl xvfb
!pip install pyvirtualdisplay
!apt install xvfb -y
!pip install piglet
!pip3 install box2d-py
!pip3 install gym[Box_2D]
!pip install tensorflow==2.3.1 gym keras-rl2
```

```
Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
ERROR: Could not find a version that satisfies the requirement python-
opengl (from versions: none)
ERROR: No matching distribution found for python-opengl
Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
Collecting pyvirtualdisplay
  Downloading PyVirtualDisplay-3.0-py3-none-any.whl (15 kB)
Installing collected packages: pyvirtualdisplay
Successfully installed pyvirtualdisplay-3.0
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following package was automatically installed and is no longer
required:
  libnvidia-common-460
Use 'apt autoremove' to remove it.
The following NEW packages will be installed:
  xvfb
0 upgraded, 1 newly installed, 0 to remove and 20 not upgraded.
Need to get 785 kB of archives.
After this operation, 2,271 kB of additional disk space will be used.
Get:1 http://archive.ubuntu.com/ubuntu bionic-updates/universe amd64
```

```
xvfb amd64 2:1.19.6-1ubuntu4.11 [785 kB]
Fetched 785 kB in 0s (2,637 kB/s)
Selecting previously unselected package xvfb.
(Reading database ... 159447 files and directories currently
installed.)
Preparing to unpack .../xvfb_2%3a1.19.6-1ubuntu4.11_amd64.deb ...
Unpacking xvfb (2:1.19.6-1ubuntu4.11) ...
Setting up xvfb (2:1.19.6-1ubuntu4.11) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
Collecting piglet
  Downloading piglet-1.0.0-py2.py3-none-any.whl (2.2 kB)
Collecting piglet-templates
  Downloading piglet_templates-1.3.0-py3-none-any.whl (67 kB)
Requirement already satisfied: attrs in /usr/local/lib/python3.7/dist-packages
(from piglet-templates->piglet) (22.1.0)
Requirement already satisfied: astunparse in
/usr/local/lib/python3.7/dist-packages (from piglet-templates->piglet)
(1.6.3)
Requirement already satisfied: pyparsing in
/usr/local/lib/python3.7/dist-packages (from piglet-templates->piglet)
(3.0.9)
Requirement already satisfied: markupsafe in
/usr/local/lib/python3.7/dist-packages (from piglet-templates->piglet)
(2.0.1)
Requirement already satisfied: six<2.0,>=1.6.1 in
/usr/local/lib/python3.7/dist-packages (from astunparse->piglet-
templates->piglet) (1.15.0)
Requirement already satisfied: wheel<1.0,>=0.23.0 in
/usr/local/lib/python3.7/dist-packages (from astunparse->piglet-
templates->piglet) (0.37.1)
Installing collected packages: piglet-templates, piglet
Successfully installed piglet-1.0.0 piglet-templates-1.3.0
Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
Collecting box2d-py
  Downloading box2d_py-2.3.8-cp37-cp37m-manylinux1_x86_64.whl (448 kB)
ple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: gym[Box_2D] in
/usr/local/lib/python3.7/dist-packages (0.25.2)
WARNING: gym 0.25.2 does not provide the extra 'box_2d'
Requirement already satisfied: importlib-metadata>=4.8.0 in
/usr/local/lib/python3.7/dist-packages (from gym[Box_2D]) (4.12.0)
Requirement already satisfied: gym-notices>=0.0.4 in
/usr/local/lib/python3.7/dist-packages (from gym[Box_2D]) (0.0.8)
Requirement already satisfied: numpy>=1.18.0 in
/usr/local/lib/python3.7/dist-packages (from gym[Box_2D]) (1.21.6)
Requirement already satisfied: cloudpickle>=1.2.0 in
/usr/local/lib/python3.7/dist-packages (from gym[Box_2D]) (1.5.0)
```

```

Requirement already satisfied: zipp>=0.5 in
/usr/local/lib/python3.7/dist-packages (from importlib-
metadata>=4.8.0->gym[Box_2D]) (3.8.1)
Requirement already satisfied: typing-extensions>=3.6.4 in
/usr/local/lib/python3.7/dist-packages (from importlib-
metadata>=4.8.0->gym[Box_2D]) (4.1.1)
Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
Collecting tensorflow==2.3.1
  Downloading tensorflow-2.3.1-cp37-cp37m-manylinux2010_x86_64.whl
(320.4 MB)
Requirement already satisfied: gym in /usr/local/lib/python3.7/dist-packages
(0.25.2)
Collecting keras-rl2
  Downloading keras_rl2-1.0.5-py3-none-any.whl (52 kB)
  -manylinux1_x86_64.whl (2.9 MB)
Requirement already satisfied: absl-py>=0.7.0 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(1.2.0)
Requirement already satisfied: protobuf>=3.9.2 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(3.17.3)
Collecting tensorflow-estimator<2.4.0,>=2.3.0
  Downloading tensorflow_estimator-2.3.0-py2.py3-none-any.whl (459 kB)
Requirement already satisfied: termcolor>=1.1.0 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(1.1.0)
Requirement already satisfied: tensorboard<3,>=2.3.0 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(2.8.0)
Requirement already satisfied: wheel>=0.26 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(0.37.1)
Requirement already satisfied: six>=1.12.0 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(1.15.0)
Requirement already satisfied: grpcio>=1.8.6 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(1.48.1)
Requirement already satisfied: wrapt>=1.11.1 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(1.14.1)
Requirement already satisfied: opt-einsum>=2.3.2 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(3.3.0)
Collecting numpy<1.19.0,>=1.16.0
  Downloading numpy-1.18.5-cp37-cp37m-manylinux1_x86_64.whl (20.1 MB)
Requirement already satisfied: astunparse==1.6.3 in
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)
(1.6.3)

```

Requirement already satisfied: google-pasta>=0.1.8 in  
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)  
(0.2.0)

Requirement already satisfied: keras-preprocessing<1.2,>=1.1.1 in  
/usr/local/lib/python3.7/dist-packages (from tensorflow==2.3.1)  
(1.1.2)

Requirement already satisfied: setuptools>=41.0.0 in  
/usr/local/lib/python3.7/dist-packages (from tensorboard<3,>=2.3.0-  
>tensorflow==2.3.1) (57.4.0)

Requirement already satisfied: markdown>=2.6.8 in  
/usr/local/lib/python3.7/dist-packages (from tensorboard<3,>=2.3.0-  
>tensorflow==2.3.1) (3.4.1)

Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in  
/usr/local/lib/python3.7/dist-packages (from tensorboard<3,>=2.3.0-  
>tensorflow==2.3.1) (1.8.1)

Requirement already satisfied: werkzeug>=0.11.15 in  
/usr/local/lib/python3.7/dist-packages (from tensorboard<3,>=2.3.0-  
>tensorflow==2.3.1) (1.0.1)

Requirement already satisfied: requests<3,>=2.21.0 in  
/usr/local/lib/python3.7/dist-packages (from tensorboard<3,>=2.3.0-  
>tensorflow==2.3.1) (2.23.0)

Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in  
/usr/local/lib/python3.7/dist-packages (from tensorboard<3,>=2.3.0-  
>tensorflow==2.3.1) (0.4.6)

Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0  
in /usr/local/lib/python3.7/dist-packages (from tensorboard<3,>=2.3.0-  
>tensorflow==2.3.1) (0.6.1)

Requirement already satisfied: google-auth<3,>=1.6.3 in  
/usr/local/lib/python3.7/dist-packages (from tensorboard<3,>=2.3.0-  
>tensorflow==2.3.1) (1.35.0)

Requirement already satisfied: rsa<5,>=3.1.4 in  
/usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3-  
>tensorboard<3,>=2.3.0->tensorflow==2.3.1) (4.9)

Requirement already satisfied: pyasn1-modules>=0.2.1 in  
/usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3-  
>tensorboard<3,>=2.3.0->tensorflow==2.3.1) (0.2.8)

Requirement already satisfied: cachetools<5.0,>=2.0.0 in  
/usr/local/lib/python3.7/dist-packages (from google-auth<3,>=1.6.3-  
>tensorboard<3,>=2.3.0->tensorflow==2.3.1) (4.2.4)

Requirement already satisfied: requests-oauthlib>=0.7.0 in  
/usr/local/lib/python3.7/dist-packages (from google-auth-  
oauthlib<0.5,>=0.4.1->tensorboard<3,>=2.3.0->tensorflow==2.3.1)  
(1.3.1)

Requirement already satisfied: importlib-metadata>=4.4 in  
/usr/local/lib/python3.7/dist-packages (from markdown>=2.6.8-  
>tensorboard<3,>=2.3.0->tensorflow==2.3.1) (4.12.0)

Requirement already satisfied: typing-extensions>=3.6.4 in  
/usr/local/lib/python3.7/dist-packages (from importlib-metadata>=4.4-  
>markdown>=2.6.8->tensorboard<3,>=2.3.0->tensorflow==2.3.1) (4.1.1)

Requirement already satisfied: zipp>=0.5 in

```

/usr/local/lib/python3.7/dist-packages (from importlib-metadata>=4.4-
>markdown>=2.6.8->tensorboard<3,>=2.3.0->tensorflow==2.3.1) (3.8.1)
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in
/usr/local/lib/python3.7/dist-packages (from pyasn1-modules>=0.2.1-
>google-auth<3,>=1.6.3->tensorboard<3,>=2.3.0->tensorflow==2.3.1)
(0.4.8)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0-
>tensorboard<3,>=2.3.0->tensorflow==2.3.1) (2022.6.15)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1
in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0-
>tensorboard<3,>=2.3.0->tensorflow==2.3.1) (1.24.3)
Requirement already satisfied: chardet<4,>=3.0.2 in
/usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0-
>tensorboard<3,>=2.3.0->tensorflow==2.3.1) (3.0.4)
Requirement already satisfied: idna<3,>=2.5 in
/usr/local/lib/python3.7/dist-packages (from requests<3,>=2.21.0-
>tensorboard<3,>=2.3.0->tensorflow==2.3.1) (2.10)
Requirement already satisfied: oauthlib>=3.0.0 in
/usr/local/lib/python3.7/dist-packages (from requests-oauthlib>=0.7.0-
>google-auth-oauthlib<0.5,>=0.4.1->tensorboard<3,>=2.3.0-
>tensorflow==2.3.1) (3.2.0)
Requirement already satisfied: cloudpickle>=1.2.0 in
/usr/local/lib/python3.7/dist-packages (from gym) (1.5.0)
Requirement already satisfied: gym-notices>=0.0.4 in
/usr/local/lib/python3.7/dist-packages (from gym) (0.0.8)
Installing collected packages: numpy, tensorflow-estimator, h5py,
gast, tensorflow, keras-rl2
  Attempting uninstall: numpy
    Found existing installation: numpy 1.21.6
    Uninstalling numpy-1.21.6:
      Successfully uninstalled numpy-1.21.6
  Attempting uninstall: tensorflow-estimator
    Found existing installation: tensorflow-estimator 2.8.0
    Uninstalling tensorflow-estimator-2.8.0:
      Successfully uninstalled tensorflow-estimator-2.8.0
  Attempting uninstall: h5py
    Found existing installation: h5py 3.1.0
    Uninstalling h5py-3.1.0:
      Successfully uninstalled h5py-3.1.0
  Attempting uninstall: gast
    Found existing installation: gast 0.5.3
    Uninstalling gast-0.5.3:
      Successfully uninstalled gast-0.5.3
  Attempting uninstall: tensorflow
    Found existing installation: tensorflow
2.8.2+zzzcolab20220719082949
    Uninstalling tensorflow-2.8.2+zzzcolab20220719082949:
      Successfully uninstalled tensorflow-2.8.2+zzzcolab20220719082949
ERROR: pip's dependency resolver does not currently take into account

```

all the packages that are installed. This behaviour is the source of the following dependency conflicts.  
xarray-einstats 0.2.2 requires numpy>=1.21, but you have numpy 1.18.5 which is incompatible.  
tables 3.7.0 requires numpy>=1.19.0, but you have numpy 1.18.5 which is incompatible.  
plotnine 0.8.0 requires numpy>=1.19.0, but you have numpy 1.18.5 which is incompatible.  
jaxlib 0.3.15+cuda11.cudnn805 requires numpy>=1.19, but you have numpy 1.18.5 which is incompatible.  
jax 0.3.17 requires numpy>=1.20, but you have numpy 1.18.5 which is incompatible.  
cmdstanpy 1.0.7 requires numpy>=1.21, but you have numpy 1.18.5 which is incompatible.  
Successfully installed gast-0.3.3 h5py-2.10.0 keras-rl2-1.0.5 numpy-1.18.5 tensorflow-2.3.1 tensorflow-estimator-2.3.0

```
{"pip_warning":{"packages":["numpy"]}}
```

### Import the libraries

```
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.models import load_model
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
import numpy as np
import random
from collections import deque
```

### Create a DQN agent

#### Use the instruction below to prepare an agent

```
# Action space include 3 actions: Buy, Sell, and Sit
#Setting up the experience replay memory to deque with 1000 elements inside it
#Empty list with inventory is created that contains the stocks that were already bought
#Setting up gamma to 0.95, that helps to maximize the current reward over the long-term
#Epsilon parameter determines whether to use a random action or to use the model for the action.
#In the beginning random actions are encouraged, hence epsilon is set up to 1.0 when the model is not trained.
#And over time the epsilon is reduced to 0.01 in order to decrease the random actions and use the trained model
#We're then set the speed of decreasing epsilon in the epsilon_decay parameter
```

```
#Defining our neural network:
```

```
#Define the neural network function called _model and it just takes
```

```

the keyword self
#Define the model with Sequential()
#Define states i.e. the previous n days and stock prices of the days
#Defining 3 hidden layers in this network
#Changing the activation function to relu because mean-squared error
is used for the loss

```

## Preprocess the stock market data

The environment is given

```

import math

# prints formatted price
def formatPrice(n):
    return ("-$" if n < 0 else "$") + "{0:.2f}".format(abs(n))

# returns the vector containing stock data from a fixed file
def getStockDataVec(key):
    vec = []
    lines = open("'" + key + ".csv", "r").read().splitlines()

    for line in lines[1:]:
        vec.append(float(line.split(",")[4]))

    return vec

# returns the sigmoid
def sigmoid(x):
    return 1 / (1 + math.exp(-x))

# returns an n-day state representation ending at time t
def getState(data, t, n):
    d = t - n + 1
    block = data[d:t + 1] if d >= 0 else -d * [data[0]] + data[0:t + 1] # pad with t0
    res = []
    for i in range(n - 1):
        res.append(sigmoid(block[i + 1] - block[i]))

    return np.array([res])

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Activation
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.optimizers import Adam

```



```

from tensorflow.keras.layers import Dense
from tensorflow.keras.regularizers import l2
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.layers import Dropout

# changes
class Agent():
    def __init__(self, window_size, is_eval=False, model_name=''):

        self.nS = window_size
        self.nA = 3
        self.memory = deque([], maxlen=1000)
        self.alpha = 0.001
        self.window_size = window_size
        self.gamma = 0.95
        #Explore/Exploit
        self.epsilon = 1
        self.epsilon_min = 0.01
        self.epsilon_decay = 0.995
        # self.model = self.build_model()
        self.loss = []

        self.is_eval = is_eval
        self.model = load_model(model_name) if self.is_eval else
self.build_model()

    def build_model(self):
        # model = keras.Sequential()
        # model.add(keras.layers.Dense(24, input_dim=self.window_size,
activation='relu')) #[Input] -> Layer 1
        # # Dense: Densely connected layer
https://keras.io/layers/core/
        # # 24: Number of neurons
        # # input_dim: Number of input variables
        # # activation: Rectified Linear Unit (relu) ranges >= 0
        # model.add(keras.layers.Dense(24, activation='relu')) #Layer
2 -> 3
        # model.add(keras.layers.Dense(self.nA, activation='linear'))
#Layer 3 -> 4
        # # model.add(keras.layers.Dense(self.nA,
activation='linear')) #Layer 4 -> [output]
        # # Size has to match the output (different actions)
        # # Linear activation on the last layer
        # model.compile(loss='mean_squared_error', #Loss function:
Mean Squared Error
        #
optimizer=keras.optimizers.Adam(lr=self.alpha)) #Optimaizer: Adam
(Feel free to check other options)

        model = Sequential()

```

```

        model.add(Dense(50, input_dim=self.window_size,
kernel_initializer="uniform", kernel_regularizer=l2(0.0002),
name="LAYER____1"))
        model.add(Activation("elu"))
        model.add(BatchNormalization())
        model.add(Dropout(0.25))
        model.add(Dense(35, kernel_initializer="uniform",
kernel_regularizer=l2(0.0002), name="LAYER____2"))
        model.add(Activation("elu"))
        model.add(BatchNormalization())
        model.add(Dropout(0.25))
        model.add(Dense(20, kernel_initializer="uniform",
kernel_regularizer=l2(0.0002), name="LAYER____3"))
        model.add(Activation("elu"))
        model.add(BatchNormalization())
        model.add(Dropout(0.25))
        model.add(Dense(self.nA, activation='linear'))

    opt = Adam(lr=0.0001, beta_1=0.5, decay=0.0002)
    model.compile(loss="mean_squared_error", optimizer=opt)

    return model

def act(self, state):#act
    if np.random.rand() <= self.epsilon:
        return random.randrange(3) #Explore
    action_vals = self.model.predict(state) #Exploit: Use the NN
to predict the correct action from this state
    return np.argmax(action_vals[0])

def test_action(self, state): #Exploit
    action_vals = self.model.predict(state)
    return np.argmax(action_vals[0])

def store(self, state, action, reward, nstate, done):
    #Store the experience in memory
    self.memory.append( (state, action, reward, nstate, done) )

def expReplay(self, batch_size): ## training the neural network
    #Execute the experience replay
    minibatch = random.sample( self.memory, batch_size ) #Randomly
sample from memory

    #Convert to numpy for speed by vectorization
    x = []
    y = []
    np_array = np.array(minibatch)
    st = np.zeros((0, self.nS)) #States
    nst = np.zeros( (0, self.nS) )#Next States

```

```

        for i in range(len(np_array)): #Creating the state and next
state np arrays
            st = np.append( st, np_array[i,0], axis=0)
            nst = np.append( nst, np_array[i,3], axis=0)
            st_predict = self.model.predict(st) #Here is the speedup! I
can predict on the ENTIRE batch
            nst_predict = self.model.predict(nst)
            index = 0
            for state, action, reward, nstate, done in minibatch:
                x.append(state)
                #Predict from state
                nst_action_predict_model = nst_predict[index]
                if done == True: #Terminal: Just assign reward much like
{* (not done) - QB[state][action]}
                    target = reward
                else: #Non terminal
                    target = reward + self.gamma *
np.amax(nst_action_predict_model)
                target_f = st_predict[index]
                target_f[action] = target
                y.append(target_f)
                index += 1
            #Reshape for Keras Fit
            x_reshape = np.array(x).reshape(batch_size,self.nS)
            y_reshape = np.array(y)
            epoch_count = 1 #Epochs is the number or iterations
            hist = self.model.fit(x_reshape, y_reshape,
epochs=epoch_count, verbose=0)
            #Graph Losses
            for i in range(epoch_count):
                self.loss.append( hist.history['loss'][i] )
            #Decay Epsilon
            if self.epsilon > self.epsilon_min:
                self.epsilon *= self.epsilon_decay

```

### Train and build the model

```

import sys

if len(sys.argv) != 4:
    print ("Usage: python train.py [stock] [window] [episodes]")
    exit()

stock_name = input("Enter stock_name, window_size, Episode_count")
#Fill the given information when prompted:
#Enter stock_name = GSPC_Training_Dataset
#window_size = 10
#Episode_count = 100 or it can be 10 or 20 or 30 and so on.

window_size = input()

```

```

episode_count = input()
stock_name = str(stock_name)
window_size = int(window_size)
episode_count = int(episode_count)

agent = Agent(window_size)
data = getStockDataVec(stock_name)
l = len(data) - 1
batch_size = 32

for e in range(episode_count + 1):
    print ("Episode " + str(e) + "/" + str(episode_count))
    state = getState(data, 0, window_size + 1)

    total_profit = 0
    agent.inventory = []

    for t in range(l):
        action = agent.act(state)

        # sit
        next_state = getState(data, t + 1, window_size + 1)
        reward = 0

        if action == 1: # buy
            agent.inventory.append(data[t])
            # print ("Buy: " + formatPrice(data[t]))

        elif action == 2 and len(agent.inventory) > 0: # sell
            bought_price = agent.inventory.pop(0)
            reward = max(data[t] - bought_price, 0)
            total_profit += data[t] - bought_price
            # print ("Sell: " + formatPrice(data[t]) + " | Profit:
" + formatPrice(data[t] - bought_price))

        done = True if t == l - 1 else False
        agent.memory.append((state, action, reward, next_state,
done))
        state = next_state

    if done:
        print ("-----")
        print ("-----Episode: {} -----".format(e))
        print ("Total Profit: " + formatPrice(total_profit))

    if len(agent.memory) > batch_size:
        agent.expReplay(batch_size)

```

```

# # if e % 10 == 0:
# if e % 10 == 0:
#     agent.model.save("model_ep" + str(e))
# agent.model.save("model_ep" + str(e))

#Fill the given information when prompted:
#Enter stock_name = GSPC_Training_Dataset
#window_size = 10
#Episode_count = 100 or it can be 10 or 20 or 30 and so on.

Usage: python train.py [stock] [window] [episodes]
Enter stock_name, window_size, Episode_countGSPC_Training_Dataset
10
30
Episode 0/30
-----
-----Episode: 0 -----
Total Profit: $865.92

WARNING:tensorflow:From
/usr/local/lib/python3.7/dist-packages/tensorflow/python/training/
tracking/tracking.py:111: Model.state_updates (from
tensorflow.python.keras.engine.training) is deprecated and will be
removed in a future version.
Instructions for updating:
This property should not be used in TensorFlow 2.0, as updates are
applied automatically.
WARNING:tensorflow:From
/usr/local/lib/python3.7/dist-packages/tensorflow/python/training/
tracking/tracking.py:111: Layer.updates (from
tensorflow.python.keras.engine.base_layer) is deprecated and will be
removed in a future version.
Instructions for updating:
This property should not be used in TensorFlow 2.0, as updates are
applied automatically.

Episode 1/30
-----
-----Episode: 1 -----
Total Profit: $7141.67
Episode 2/30
-----
-----Episode: 2 -----
Total Profit: $6871.84
Episode 3/30
-----
-----Episode: 3 -----
Total Profit: $6425.62
Episode 4/30
-----
-----Episode: 4 -----

```

Total Profit: \$7306.68

Episode 5/30

-----Episode: 5 -----

Total Profit: \$7161.61

Episode 6/30

-----Episode: 6 -----

Total Profit: \$6990.42

Episode 7/30

-----Episode: 7 -----

Total Profit: \$7420.58

Episode 8/30

-----Episode: 8 -----

Total Profit: \$5741.95

Episode 9/30

-----Episode: 9 -----

Total Profit: \$7061.20

Episode 10/30

-----Episode: 10 -----

Total Profit: \$5241.04

Episode 11/30

-----Episode: 11 -----

Total Profit: \$7092.10

Episode 12/30

-----Episode: 12 -----

Total Profit: \$6816.25

Episode 13/30

-----Episode: 13 -----

Total Profit: \$6431.38

Episode 14/30

-----Episode: 14 -----

Total Profit: \$6240.14

Episode 15/30

-----Episode: 15 -----

Total Profit: \$5827.87

Episode 16/30

-----Episode: 16 -----

Total Profit: \$5912.93

Episode 17/30

```
-----
-----Episode: 17 -----
Total Profit: $4592.13
Episode 18/30
-----
-----Episode: 18 -----
Total Profit: $6146.69
Episode 19/30
-----
-----Episode: 19 -----
Total Profit: $6504.60
Episode 20/30
-----
-----Episode: 20 -----
Total Profit: $5526.26
Episode 21/30
-----
-----Episode: 21 -----
Total Profit: $6493.13
Episode 22/30
-----
-----Episode: 22 -----
Total Profit: $6761.81
Episode 23/30
-----
-----Episode: 23 -----
Total Profit: $6308.82
Episode 24/30
-----
-----Episode: 24 -----
Total Profit: $6131.78
Episode 25/30
-----
-----Episode: 25 -----
Total Profit: $6599.88
Episode 26/30
-----
-----Episode: 26 -----
Total Profit: $5144.48
Episode 27/30
-----
-----Episode: 27 -----
Total Profit: $5699.61
Episode 28/30
-----
-----Episode: 28 -----
Total Profit: $5480.21
Episode 29/30
-----
-----Episode: 29 -----
```

Total Profit: \$5844.89

Episode 30/30

-----

-----Episode: 30 -----

Total Profit: \$6726.90

*# save all the model to my google drive*

```
!scp -r /content/model_ep*
/content/drive/MyDrive/datasets/simplilearn_RL_stock_trading/model_file
```

### Evaluate the model and agent

```
import sys
```

```
from tensorflow.keras.models import load_model
```

```
if len(sys.argv) != 3:
```

```
    print ("Usage: python evaluate.py [stock] [model]")
```

```
    exit()
```

```
stock_name = input("Enter Stock_name, Model_name")
```

```
model_name = input()
```

*#Note:*

*#Fill the given information when prompted:*

*#Enter stock\_name = GSPC\_Evaluation\_Dataset*

*#Model\_name = respective model name*

```
model = load_model("" + model_name)
```

```
window_size = model.layers[0].input.shape.as_list()[1]
```

```
agent = Agent(window_size, True, model_name)
```

```
data = getStockDataVec(stock_name)
```

```
l = len(data) - 1
```

```
batch_size = 32
```

```
state = getState(data, 0, window_size + 1)
```

```
total_profit = 0
```

```
agent.inventory = []
```

```
for t in range(l):
```

```
    action = agent.act(state)
```

*# sit*

```
    next_state = getState(data, t + 1, window_size + 1)
```

```
    reward = 0
```

```
    if action == 1: # buy
```

```
        agent.inventory.append(data[t])
```



```

        print ("Buy: " + formatPrice(data[t]))

    elif action == 2 and len(agent.inventory) > 0: # sell
        bought_price = agent.inventory.pop(0)
        reward = max(data[t] - bought_price, 0)
        total_profit += data[t] - bought_price
        print ("Sell: " + formatPrice(data[t]) + " | Profit: " +
formatPrice(data[t] - bought_price))

    done = True if t == l - 1 else False
    agent.memory.append((state, action, reward, next_state, done))
    state = next_state

    if done:
        print ("-----")
        # print      ("-----Episode: {} -----".format(e))
        print (stock_name + " Total Profit: " +
formatPrice(total_profit))

        # if len(agent.memory) > batch_size:
        #     agent.expReplay(batch_size)

```

```

# GSPC_Evaluation_Dataset
# /content/model_ep30

```

```

Enter Stock_name, Model_nameGSPC_Evaluation_Dataset
/content/model_ep30

```

```

Buy: $1271.87
Buy: $1276.56
Buy: $1273.85
Buy: $1271.50
Buy: $1269.75
Buy: $1285.96
Sell: $1293.24 | Profit: $21.37
Buy: $1295.02
Sell: $1281.92 | Profit: $5.36
Sell: $1280.26 | Profit: $6.41
Buy: $1283.35
Buy: $1291.18
Buy: $1296.63
Sell: $1276.34 | Profit: $4.84
Sell: $1286.12 | Profit: $16.37
Sell: $1307.10 | Profit: $21.14
Buy: $1310.87
Buy: $1319.05
Sell: $1324.57 | Profit: $29.55
Buy: $1320.88
Sell: $1321.87 | Profit: $38.52
Buy: $1329.15
Sell: $1328.01 | Profit: $36.83

```

Sell: \$1343.01		Profit: \$46.38
Sell: \$1315.44		Profit: \$4.57
Sell: \$1307.40		Profit: -\$11.65
Buy: \$1306.10		
Buy: \$1319.88		
Sell: \$1327.22		Profit: \$6.34
Sell: \$1306.33		Profit: -\$22.82
Sell: \$1330.97		Profit: \$24.87
Sell: \$1321.15		Profit: \$1.27
Buy: \$1310.13		
Buy: \$1321.82		
Buy: \$1320.02		
Buy: \$1295.11		
Buy: \$1304.28		
Sell: \$1296.39		Profit: -\$13.74
Buy: \$1281.87		
Sell: \$1256.88		Profit: -\$64.94
Buy: \$1298.38		
Sell: \$1293.77		Profit: -\$26.25
Sell: \$1297.54		Profit: \$2.43
Buy: \$1319.44		
Sell: \$1328.26		Profit: \$23.98
Buy: \$1325.83		
Sell: \$1332.41		Profit: \$50.54
Buy: \$1332.87		
Sell: \$1335.54		Profit: \$37.16
Buy: \$1333.51		
Buy: \$1324.46		
Sell: \$1314.52		Profit: -\$4.92
Sell: \$1319.68		Profit: -\$6.15
Sell: \$1305.14		Profit: -\$27.73
Buy: \$1312.62		
Sell: \$1330.36		Profit: -\$3.15
Buy: \$1337.38		
Sell: \$1347.24		Profit: \$22.78
Buy: \$1355.66		
Buy: \$1360.48		
Buy: \$1363.61		
Sell: \$1361.22		Profit: \$48.60
Buy: \$1356.62		
Sell: \$1347.32		Profit: \$9.94
Sell: \$1340.20		Profit: -\$15.46
Buy: \$1346.29		
Buy: \$1357.16		
Sell: \$1342.08		Profit: -\$18.40
Sell: \$1328.98		Profit: -\$34.63
Sell: \$1340.68		Profit: -\$15.94
Buy: \$1343.60		
Buy: \$1333.27		
Buy: \$1317.37		

Sell: \$1320.47 | Profit: -\$25.82  
Buy: \$1331.10  
Sell: \$1286.17 | Profit: -\$70.99  
Buy: \$1279.56  
Sell: \$1271.83 | Profit: -\$71.77  
Buy: \$1287.87  
Buy: \$1267.64  
Buy: \$1271.50  
Buy: \$1295.52  
Sell: \$1287.14 | Profit: -\$46.13  
Buy: \$1296.67  
Buy: \$1320.64  
Sell: \$1339.67 | Profit: \$22.30  
Buy: \$1337.88  
Sell: \$1339.22 | Profit: \$8.12  
Buy: \$1353.22  
Sell: \$1319.49 | Profit: \$39.93  
Sell: \$1313.64 | Profit: \$25.77  
Sell: \$1317.72 | Profit: \$50.08  
Sell: \$1308.87 | Profit: \$37.37  
Buy: \$1316.14  
Sell: \$1326.73 | Profit: \$31.21  
Sell: \$1325.84 | Profit: \$29.17  
Sell: \$1343.80 | Profit: \$23.16  
Sell: \$1345.02 | Profit: \$7.14  
Sell: \$1337.43 | Profit: -\$15.79  
Buy: \$1331.94  
Sell: \$1304.89 | Profit: -\$11.25  
Sell: \$1300.67 | Profit: -\$31.27  
Buy: \$1286.94  
Buy: \$1260.34  
Sell: \$1199.38 | Profit: -\$87.56  
Sell: \$1172.53 | Profit: -\$87.81  
Buy: \$1178.81  
Buy: \$1193.89  
Buy: \$1140.65  
Buy: \$1123.82  
Sell: \$1162.35 | Profit: -\$16.46  
Sell: \$1159.27 | Profit: -\$34.62  
Sell: \$1176.80 | Profit: \$36.15  
Sell: \$1210.08 | Profit: \$86.26  
Buy: \$1204.42  
Sell: \$1165.24 | Profit: -\$39.18  
Buy: \$1198.62  
Buy: \$1185.90  
Buy: \$1172.87  
Buy: \$1209.11  
Buy: \$1204.09  
Sell: \$1202.09 | Profit: \$3.47  
Sell: \$1129.56 | Profit: -\$56.34

Sell: \$1136.43 | Profit: -\$36.44  
 Buy: \$1175.38  
 Buy: \$1151.06  
 Buy: \$1160.40  
 Buy: \$1131.42  
 Sell: \$1123.95 | Profit: -\$85.16  
 Buy: \$1144.03  
 Buy: \$1155.46  
 Sell: \$1207.25 | Profit: \$3.16  
 Sell: \$1203.66 | Profit: \$28.28  
 Buy: \$1224.58  
 Sell: \$1209.88 | Profit: \$58.82  
 Buy: \$1215.39  
 Sell: \$1254.19 | Profit: \$93.79  
 Buy: \$1229.05  
 Buy: \$1242.00  
 Sell: \$1284.59 | Profit: \$153.17  
 Buy: \$1253.30  
 Buy: \$1237.90  
 Buy: \$1261.15  
 Sell: \$1253.23 | Profit: \$109.20  
 Sell: \$1261.12 | Profit: \$105.66  
 Sell: \$1229.10 | Profit: \$4.52  
 Sell: \$1251.78 | Profit: \$36.39  
 Sell: \$1257.81 | Profit: \$28.76  
 Buy: \$1216.13  
 Buy: \$1192.98  
 Sell: \$1188.04 | Profit: -\$53.96  
 Sell: \$1161.79 | Profit: -\$91.51  
 Sell: \$1158.67 | Profit: -\$79.23  
 Buy: \$1195.19  
 Sell: \$1246.96 | Profit: -\$14.19  
 Sell: \$1244.58 | Profit: \$28.45  
 Sell: \$1261.01 | Profit: \$68.03  
 Buy: \$1234.35  
 Sell: \$1255.19 | Profit: \$60.00  
 Buy: \$1236.47  
 Sell: \$1225.73 | Profit: -\$8.62  
 Sell: \$1211.82 | Profit: -\$24.65  
 Buy: \$1215.75  
 Buy: \$1219.66  
 Sell: \$1243.72 | Profit: \$27.97  
 Sell: \$1254.00 | Profit: \$34.34

-----  
 GSPC\_Evaluation\_Dataset Total Profit: \$445.39

**Note: Run the training section for considerable episodes so that while evaluating the model it can generate significant profit.**